

Learning to Hack on Postgres Planner

Melanie Plageman

Goals

- Provide a tangible, trivial example of adding a fix to PostgreSQL planner
- Start a discussion on specifying where to add new optimizations to PostgreSQL planner

Table of Contents

- Postgres Planner Basics
 - Query Planning
- Guidelines for New Optimizations
- Case Study:
 - Current Plan and Semantics
 - Identifying a Target Plan and Query Tree Transformation
 - Constant Folding
 - ANY Sublink Pullup
- Resources and Discussion

Query Planning

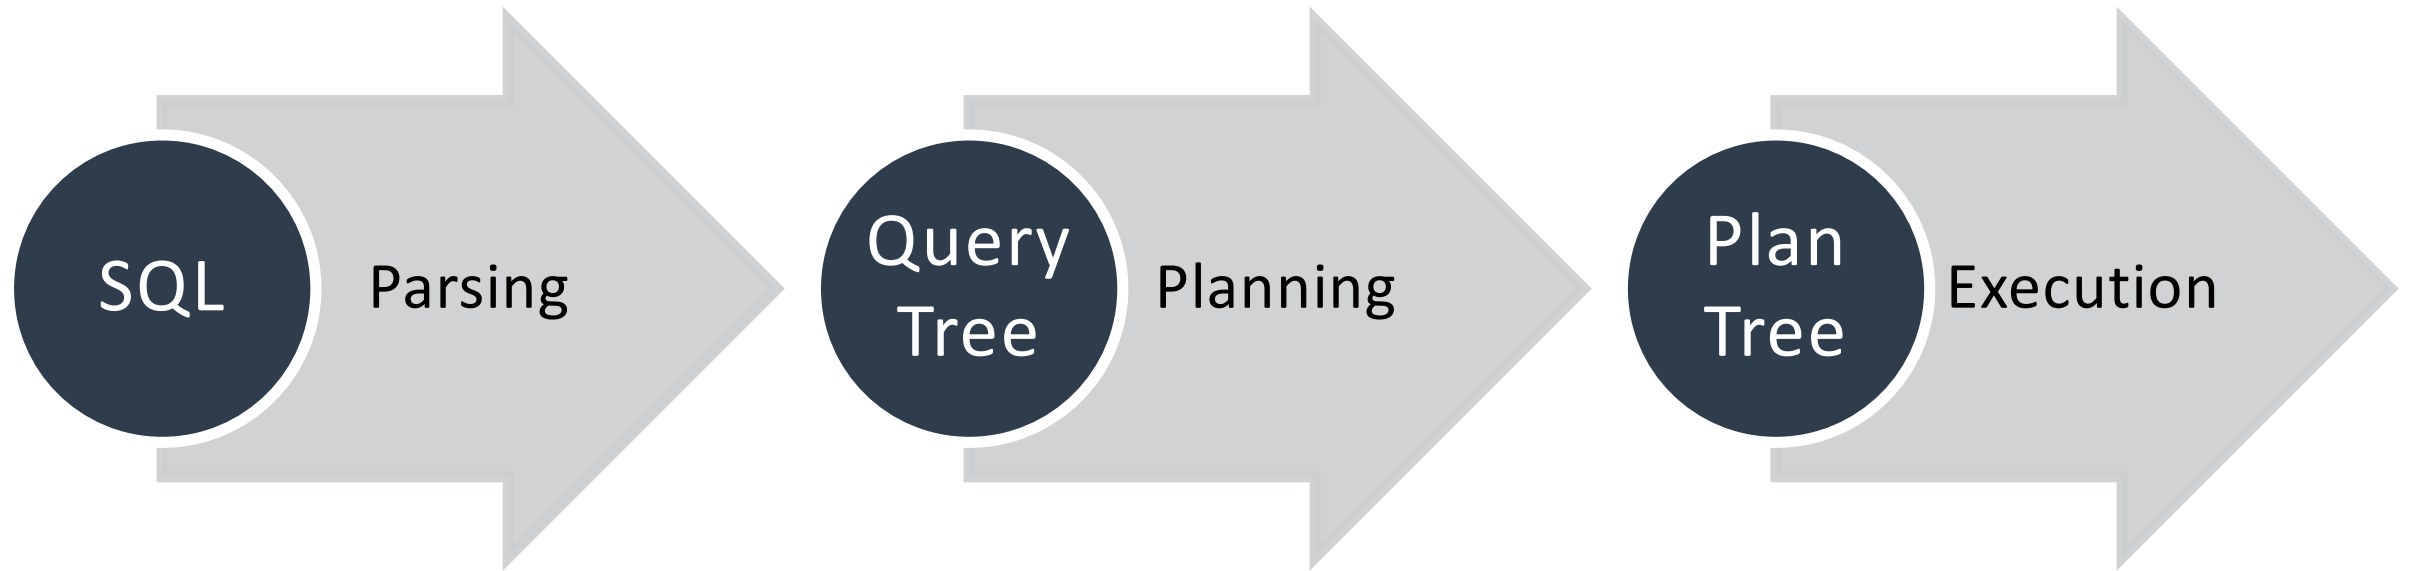
SQL statement to plan tree

```
# SELECT a FROM foo;
```



a
1
2
4

(3 rows)




Query Tree

```
# SET debug_print_parse  
TO on;
```

```
# SELECT a  
FROM foo  
WHERE 1 + 2 = 4;
```

```
{QUERY  
:rtable (  
  {RTE  
  :eref  
    {ALIAS  
    :aliasname foo  
    :colnames ("a")  
:jointree  
  {FROMEXPR  
  :quals  
    {OPEXPR  
    :args (  
      {OPEXPR  
      :args (  
        {CONST  
        :constvalue 4 [ 1 ... ]  
        {CONST  
        :constvalue 4 [ 2 ... ]  
        {CONST  
        :constvalue 4 [ 4 ... ]  
:targetList (  
  {TARGETENTRY  
  :expr  
    {VAR  
    :resname a
```



```

{QUERY
:rtable (
  {RTE
  :eref
    {ALIAS
    :aliasname foo
    :colnames ("a")}
  :jointree
    {FROMEXPR
    :quals
      {OEXPR
      :args (
        {OEXPR
        :args (
          {CONST
          :constvalue 4 [ 1 ... ]
          {CONST
          :constvalue 4 [ 2 ... ]
          {CONST
          :constvalue 4 [ 4 ... ]
        :targetList (
          {TARGETENTRY
          :expr
            {VAR
            :resname a
  
```


Semantic Optimization

```
# SELECT a FROM foo WHERE 1 + 2 = 4;
```

1 + 2 = 4

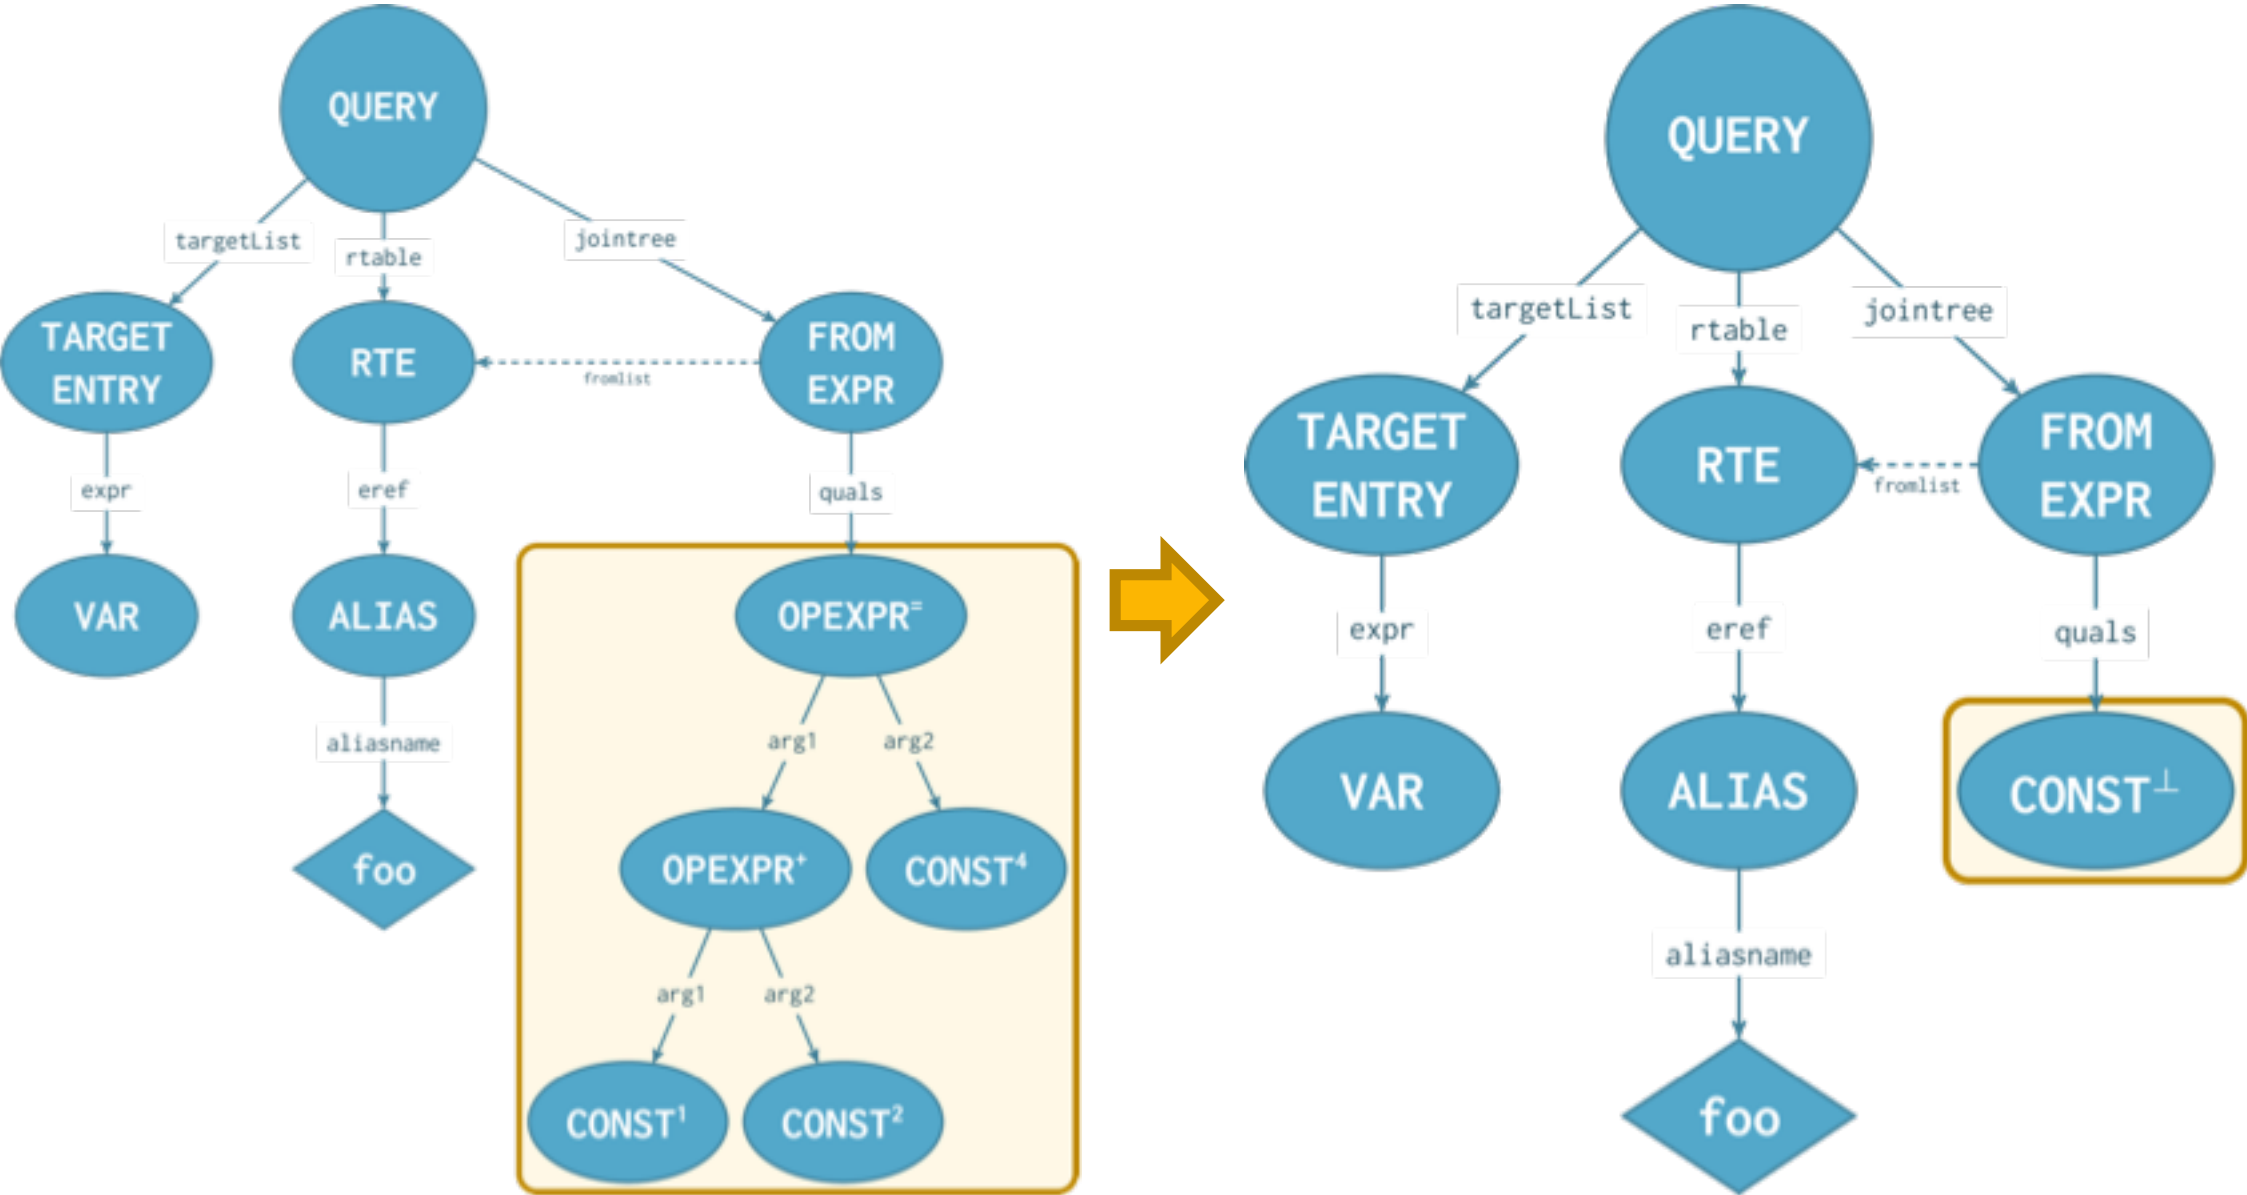


FALSE


```
# SELECT a FROM foo WHERE FALSE;
```

SELECT a FROM foo WHERE 1 + 2 = 4;

SELECT a FROM foo WHERE FALSE;



Cost-based Optimization

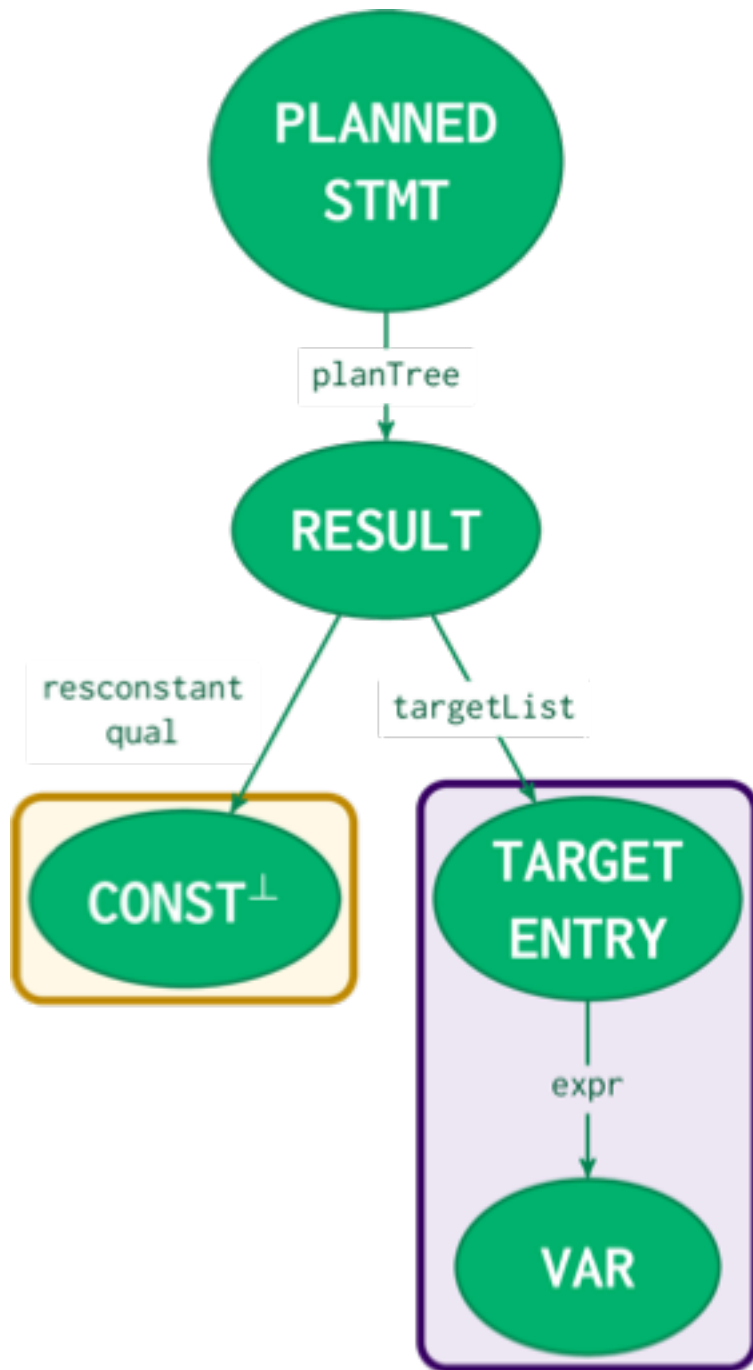


Plan Tree

```
# SET debug_print_plan  
    TO on;
```

```
# SELECT a  
    FROM foo  
    WHERE 1 + 2 = 4;
```

```
{PLANNEDSTMT  
:planTree  
  {RESULT  
    :targetlist (  
      {TARGETENTRY  
        :expr  
          {VAR  
            :resname a  
            :resconstantqual (  
              {CONST  
                :constvalue 1 [ 0 ... ]
```



```
{PLANNEDSTMT
:planTree
{RESULT
:targetlist (
{TARGETENTRY
:expr
{VAR
:resname a
:resconstantqual (
{CONST
:constvalue 1 [ 0 ... ]
```

Guidelines for New Optimizations

① Does it always retain semantic correctness?

$$\begin{array}{c} A \bowtie (B \bowtie C) \\ \neq \\ (A \bowtie B) \bowtie C \end{array}$$

② Does it inhibit downstream optimizations?

Optimization Order Matters

An optimization for one query can be a regression for another
Planning steps have expectations for the query tree

② Does it inhibit downstream optimizations?

Optimization Order Matters

SELECT * FROM A, B, C

WHERE a IN (

SELECT b FROM B WHERE b = 5

) AND a = c

AND c = 7;

c = 7

c = a \Rightarrow a = 7

{ a, c, 7 } =

b = 5

{ b, 5 } =

② Does it inhibit downstream optimizations?

Optimization Order Matters

$$c = 7$$

$$c = a \Rightarrow a = 7$$

$$\{a, c, 7\} =$$

$$b = 5$$

$$\{b, 5\} =$$

1. Pullup

2. Pre-process

$$c = 7$$

$$c = a \Rightarrow a = 7$$

$$b = 5$$

$$a = b \Rightarrow a = 5,$$

$$\Rightarrow c = 5$$

$$\{a, c, 7, b, 5\} =$$

② Does it inhibit downstream optimizations?

Optimization Order Matters

```
SELECT * FROM A, B, C
```

QUERY PLAN

```
WHERE a IN (
```

```
    SELECT b FROM B WHERE b = 5
```

Result

```
) AND a = c
```

One-Time Filter: false

```
AND c = 7;
```

② Does it inhibit downstream optimizations?

Order matters

An optimization for one query can be a regression for another

Planning steps have expectations for the query tree

② Does it inhibit downstream optimizations?


Order matters

An optimization for one query can be a regression for another

Planning steps have expectations for the query tree

③ Is the improvement in execution time worth the cost in planning time?

No in the case of **exhaustive** join order = $O(n!)$



④ Is the complexity cost commensurate with the performance benefit?

- **Narrow use cases**
- Optimizations for obscure features
- New APIs without reuse potential

Case Study

Adding a planner improvement

Table "public.foo"

Column	Type
a	integer

Table "public.bar"

Column	Type
b	integer

```
# SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

NULL \approx Unknown

p

q

p OR q

p AND q

p = q

NULL \approx Unknown

p	q	p OR q	p AND q	p = q
TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	FALSE	TRUE	FALSE	FALSE
FALSE	FALSE	FALSE	FALSE	TRUE

NULL \approx Unknown

p	q	p OR q	p AND q	p = q
TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	FALSE	TRUE	FALSE	FALSE
FALSE	FALSE	FALSE	FALSE	TRUE
TRUE	NULL	TRUE	NULL	NULL
FALSE	NULL	NULL	FALSE	NULL
NULL	NULL	NULL	NULL	NULL

EXPLAIN Output?

```
# EXPLAIN SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

QUERY PLAN

Result

One-Time Filter: false



EXPLAIN Output!

```
# EXPLAIN SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

QUERY PLAN

Result

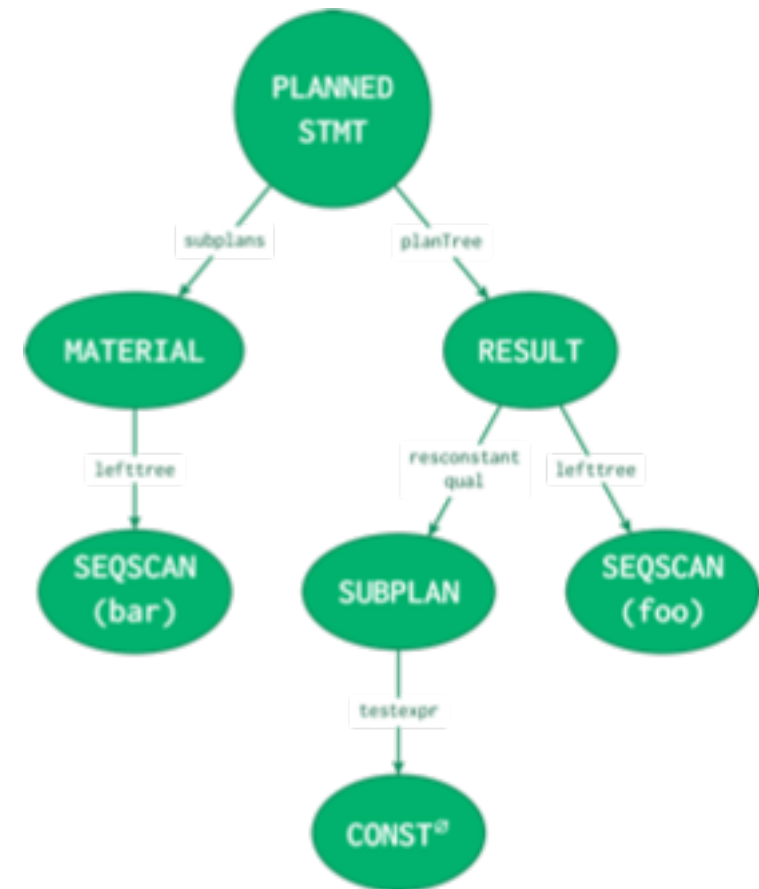
One-Time Filter: (SubPlan 1)

→ Seq Scan on foo

SubPlan 1

→ Materialize

→ Seq Scan on bar



Target Transformation

1. **Characterize the query**
2. Find analogues
3. Identify transformations

Provably **UNTRUE** quals

```
# SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

NULL = ANY(SELECT b FROM bar)



UNTRUE

```
# SELECT a FROM foo WHERE UNTRUE;
```

Target Transformation

1. Characterize the query
2. **Find analogues**
3. Identify transformations

EXPLAIN SELECT a FROM foo WHERE FALSE;

QUERY PLAN

Result

One-Time Filter: false



```
# EXPLAIN SELECT a FROM foo WHERE NULL = 7;
```



QUERY PLAN

Result

One-Time Filter: false



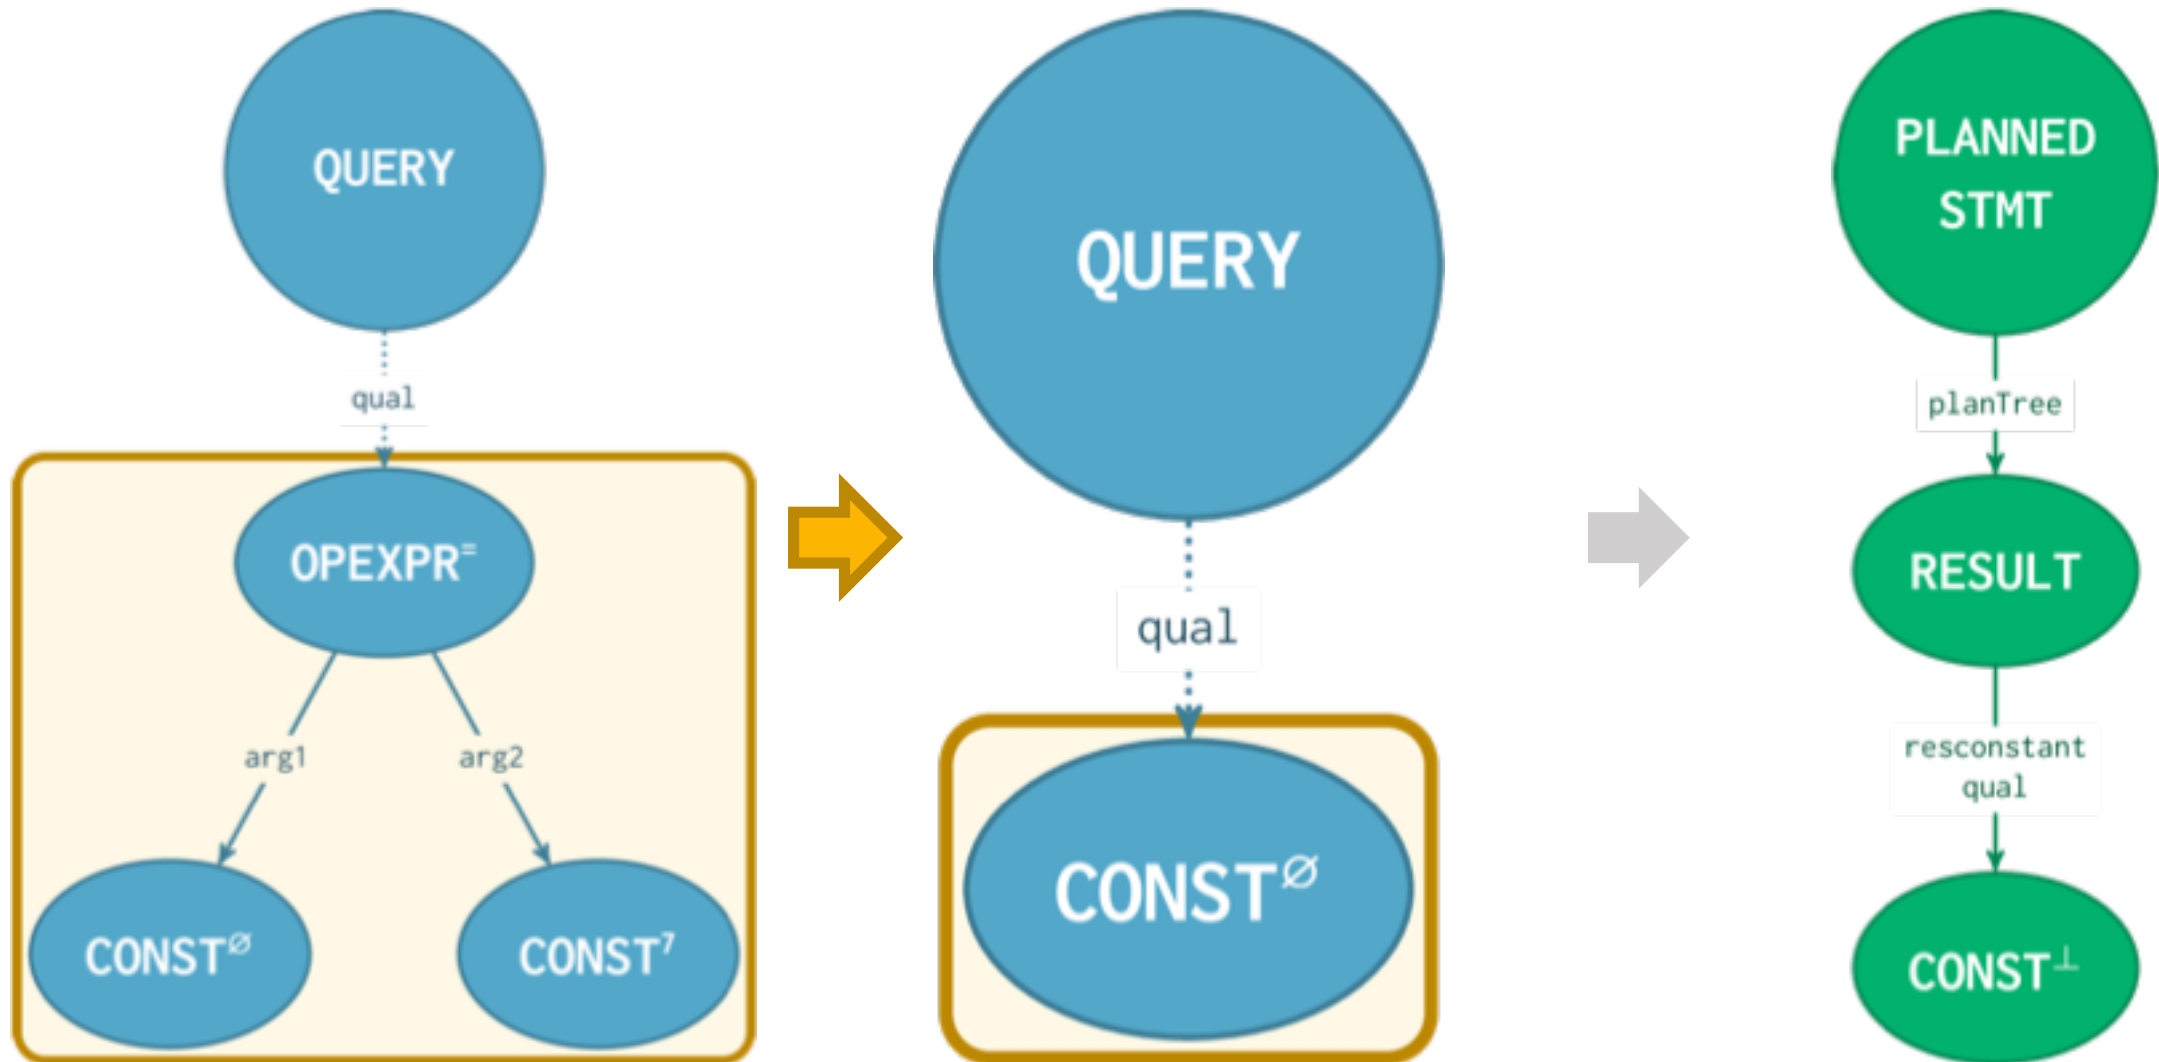
A Note on Notation



Target Transformation

1. Characterize the query
2. Find analogues
3. **Identify transformations**

SELECT a FROM foo WHERE NULL = 7;



```
# EXPLAIN SELECT a FROM foo WHERE NULL = (SELECT b FROM bar);
```

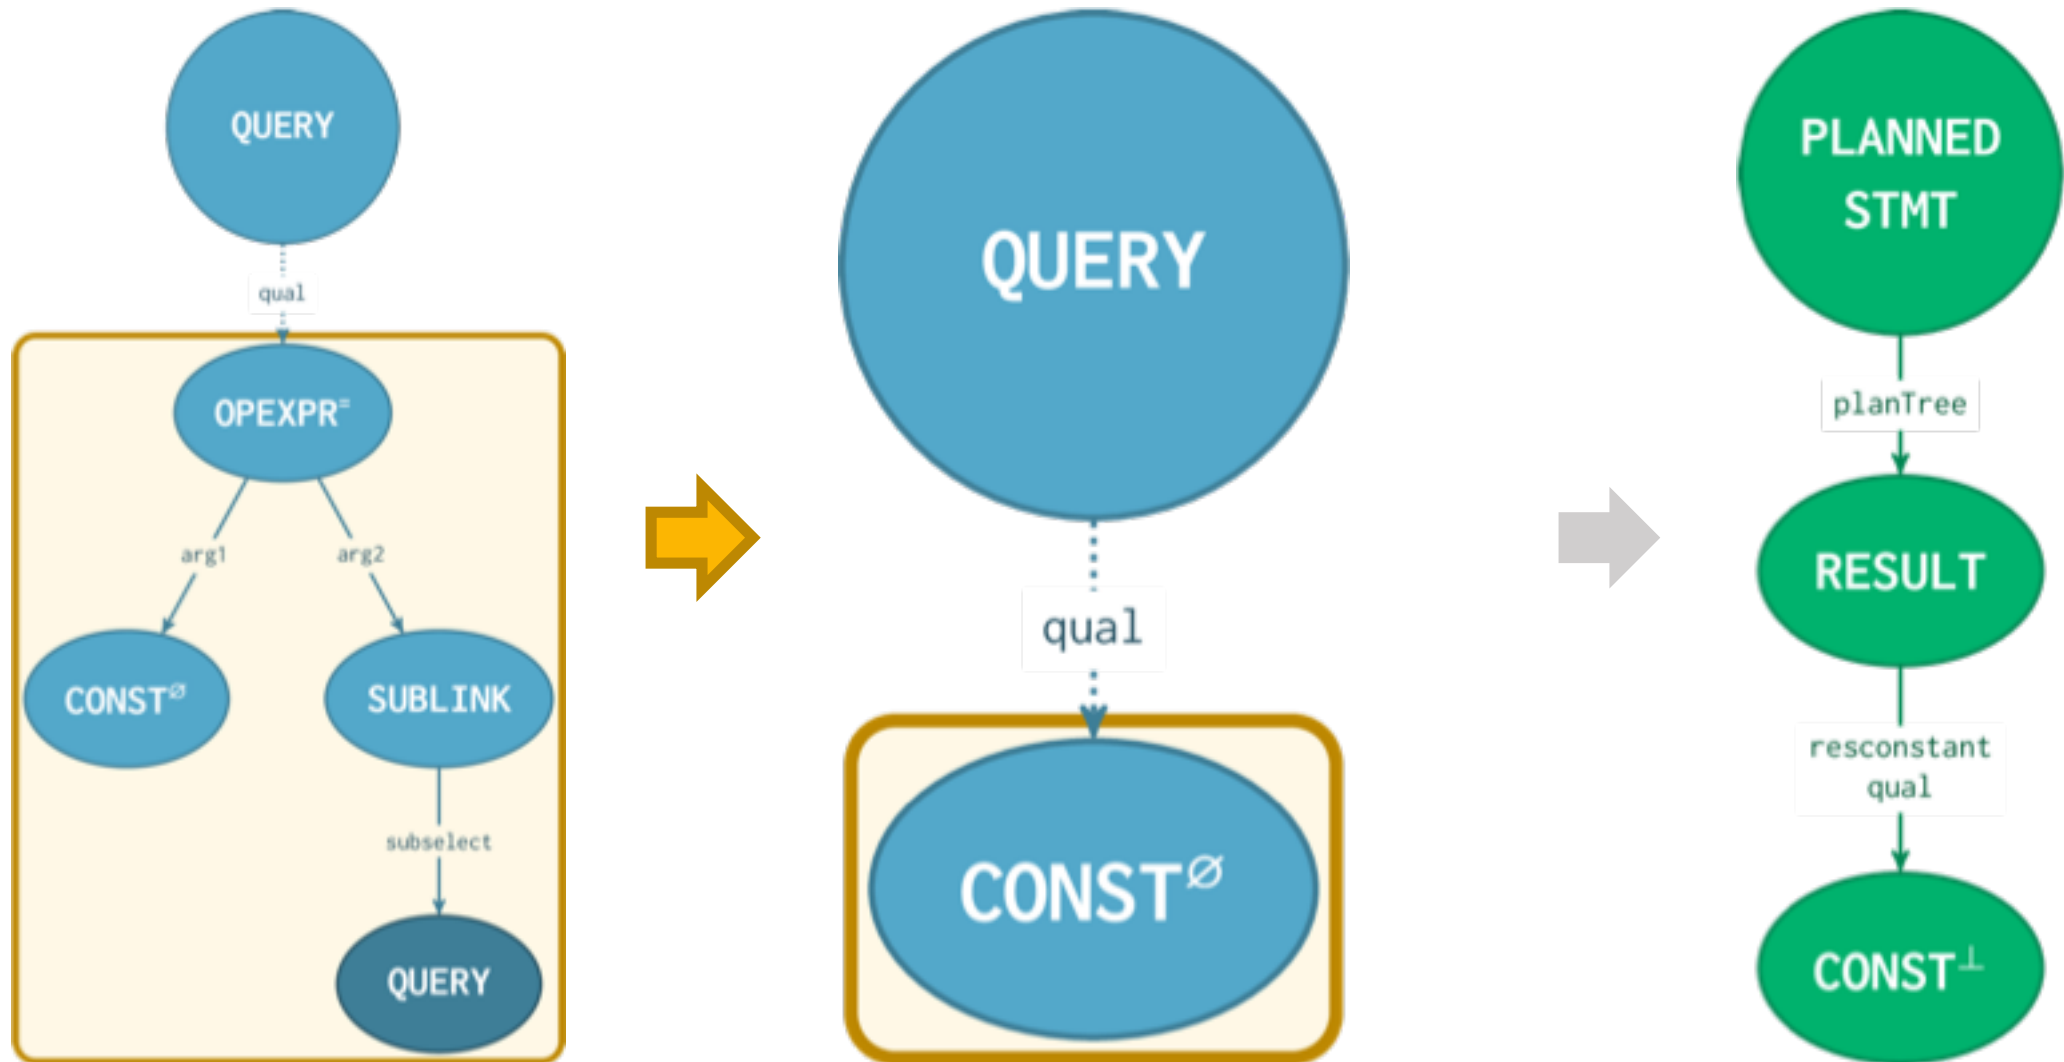
QUERY PLAN

Result

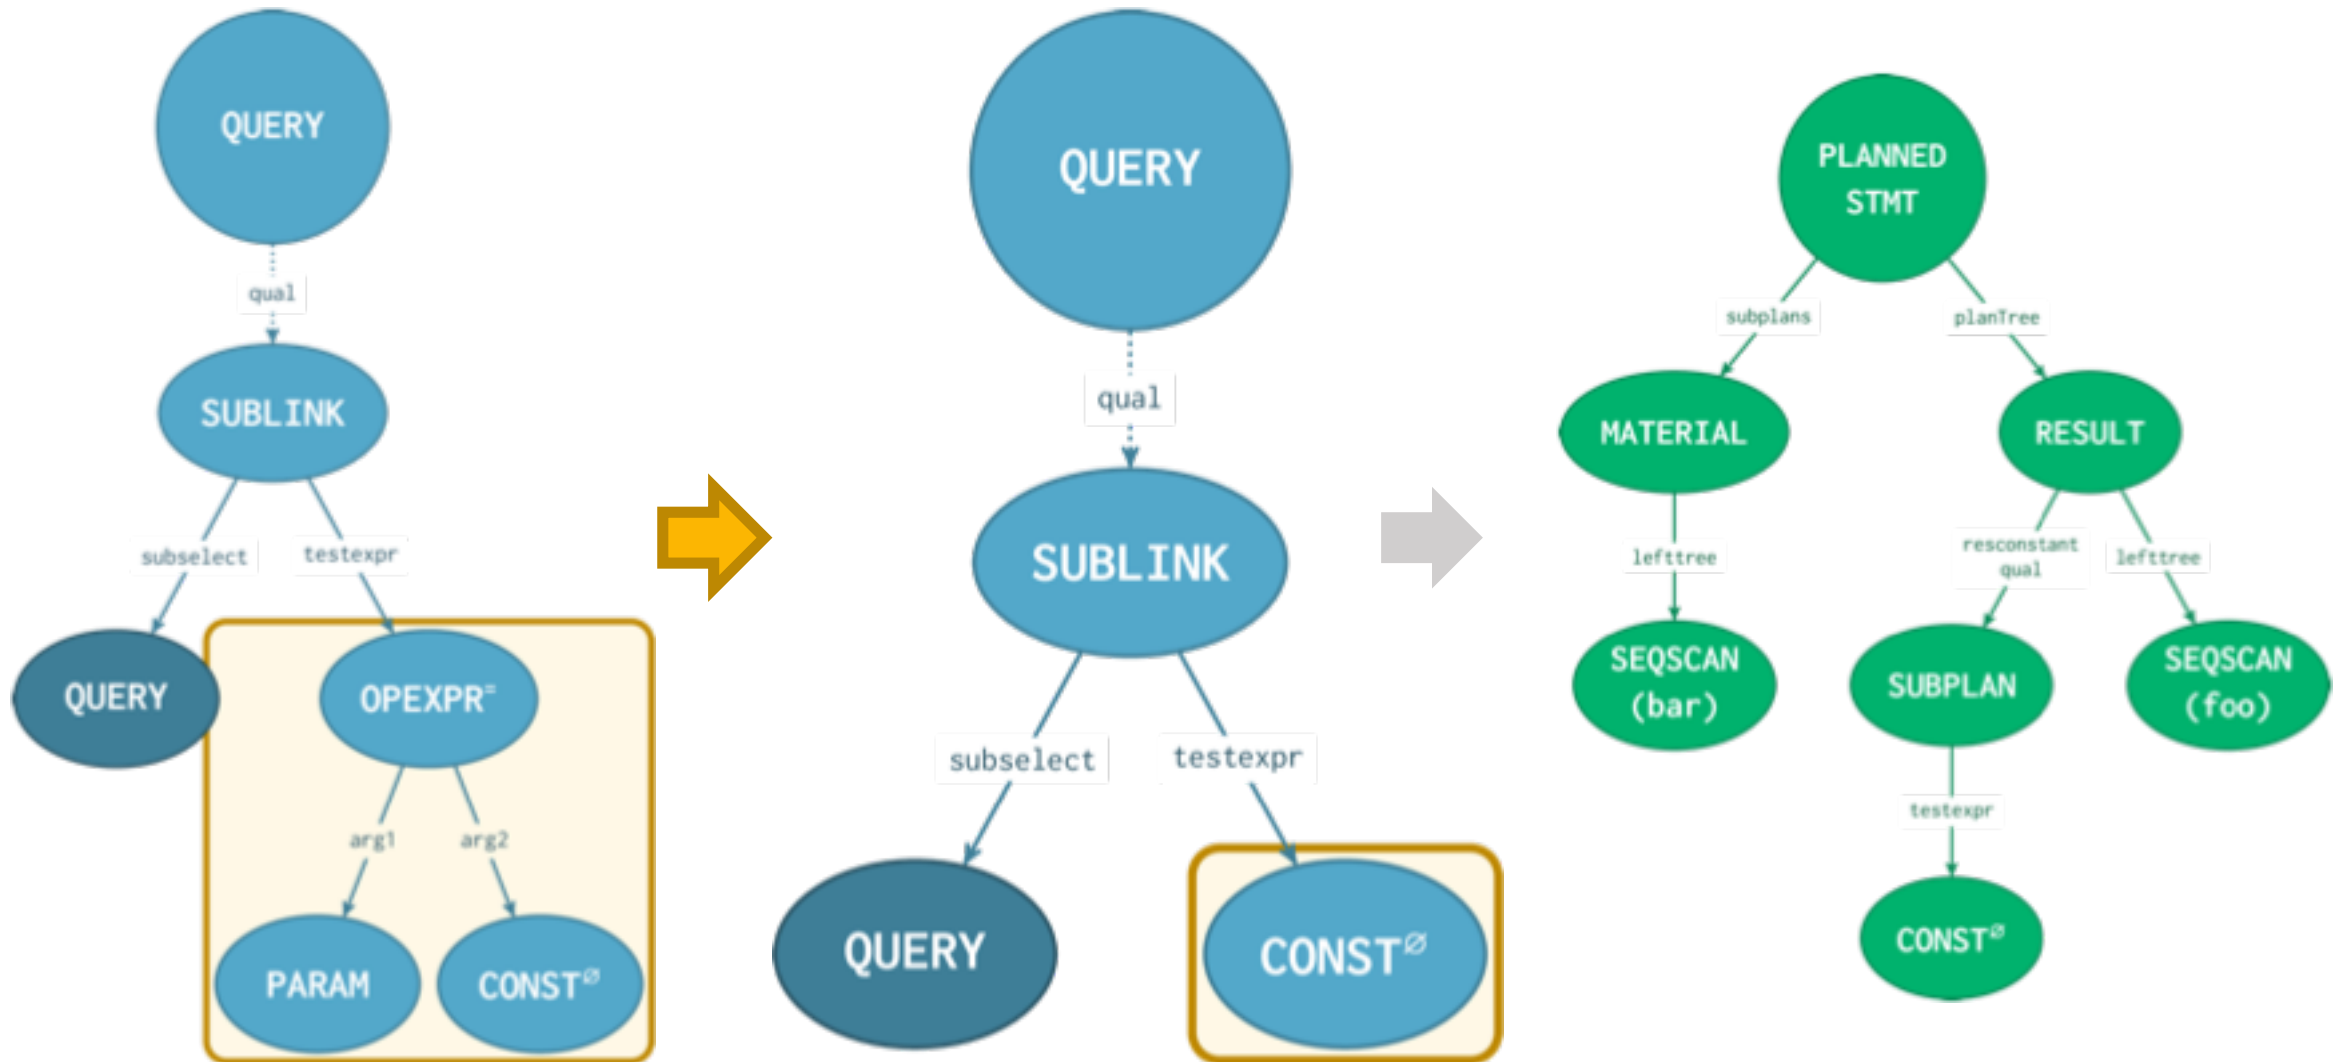
One-Time Filter: **false**



SELECT a FROM foo WHERE NULL = (SELECT b FROM bar);




SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);




SELECT *a* FROM *foo* WHERE ...


NULL = 7



NULL = (SELECT *b* FROM *bar*)




NULL = ANY(SELECT *b* FROM *bar*)




SELECT *a* FROM *foo* WHERE ...

NULL = 7

NULL = (SELECT *b* FROM *bar*)



NULL = ANY(SELECT *b* FROM *bar*)



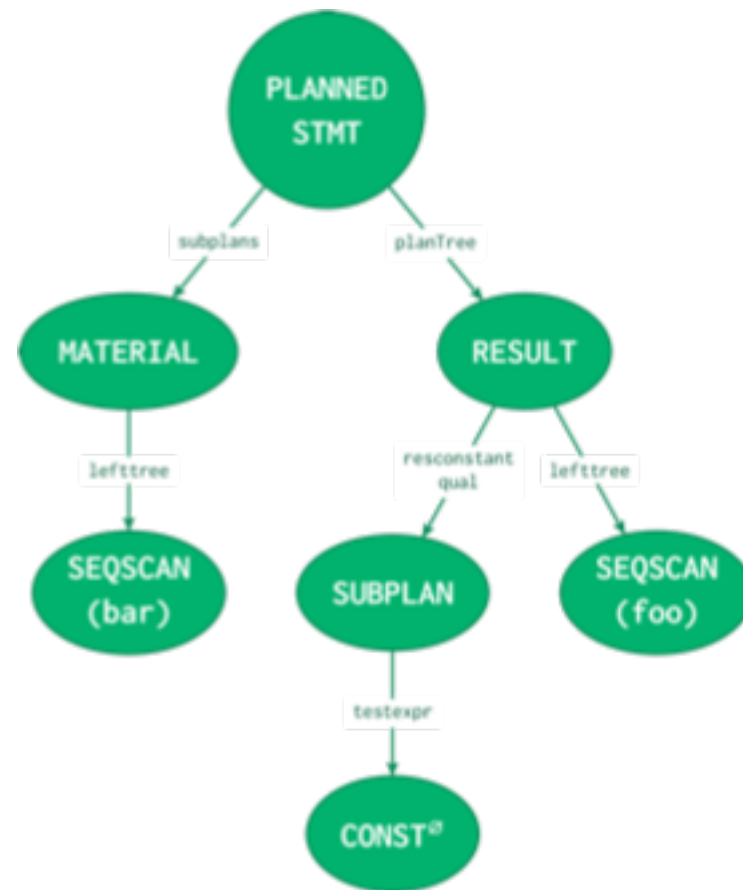
EXPLAIN SELECT a FROM foo WHERE ...

NULL = 7

NULL = (SELECT b FROM bar)



NULL = ANY(SELECT b FROM bar)



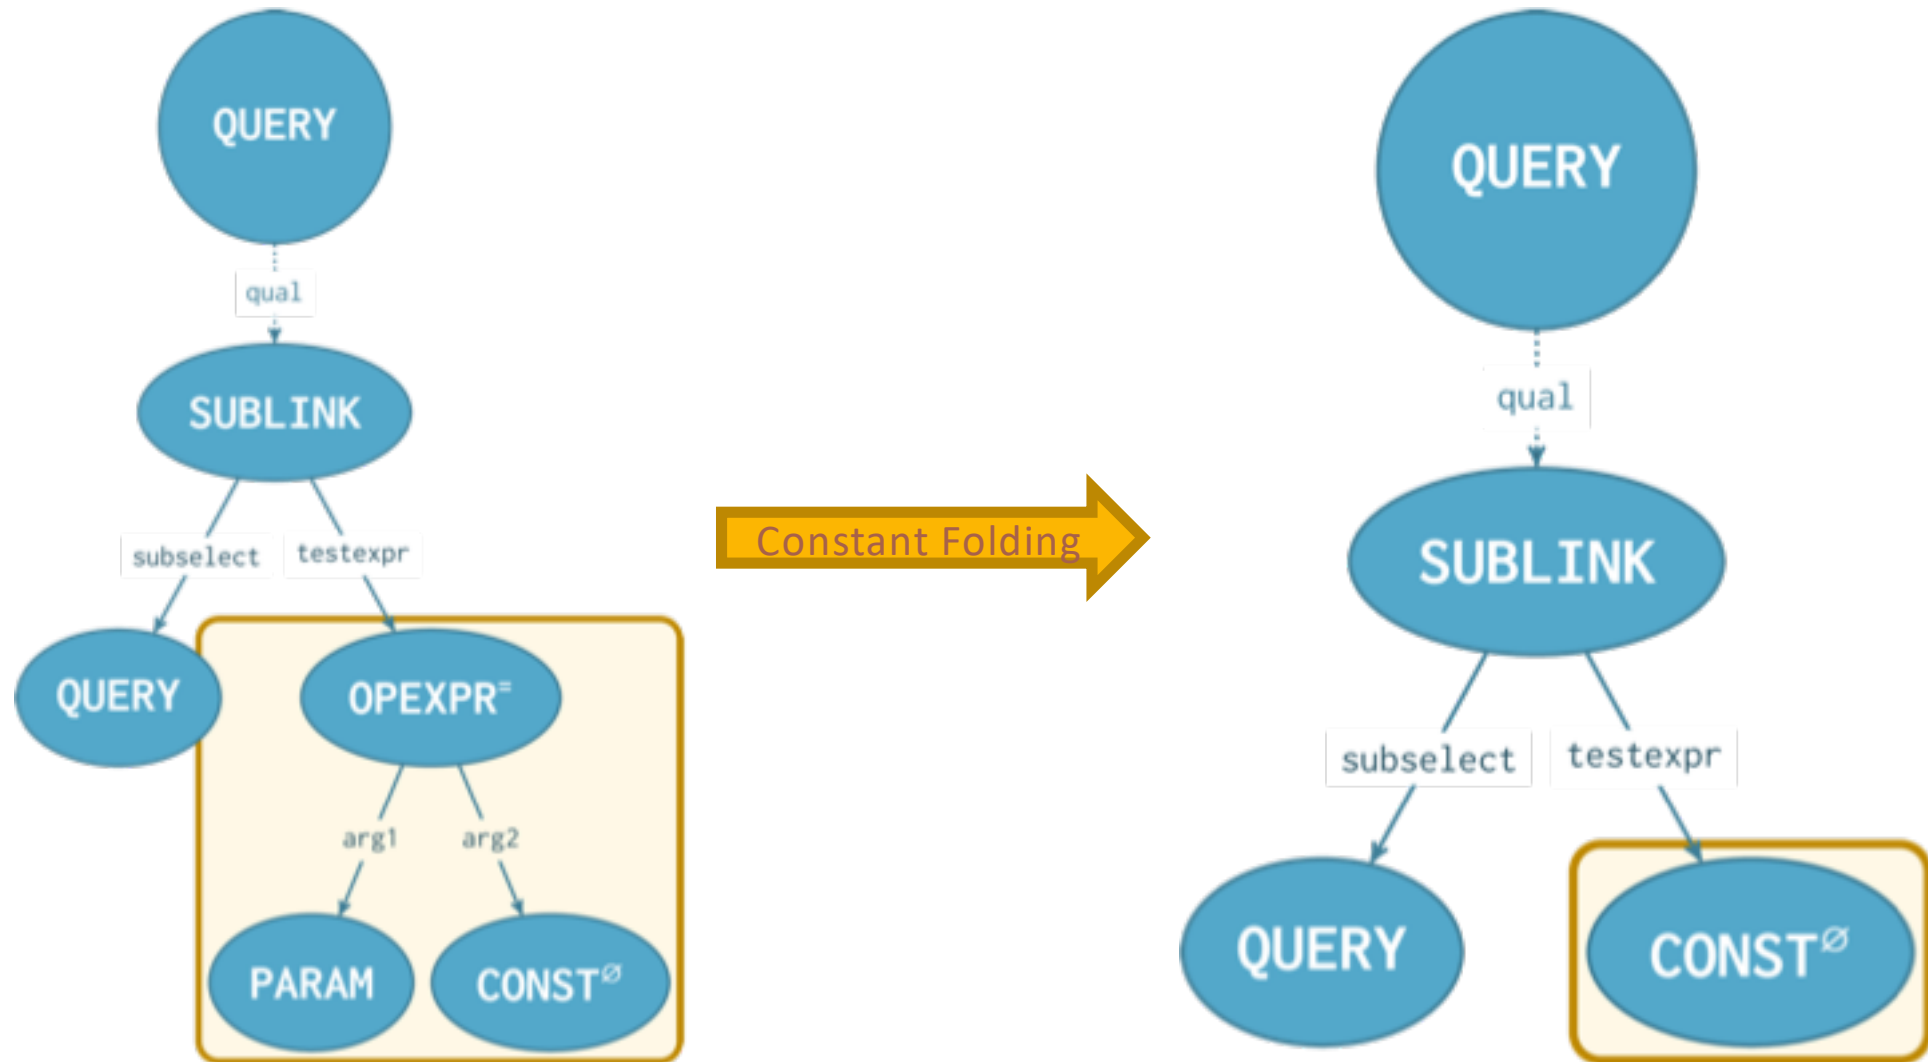
Two  s

Constant Folding

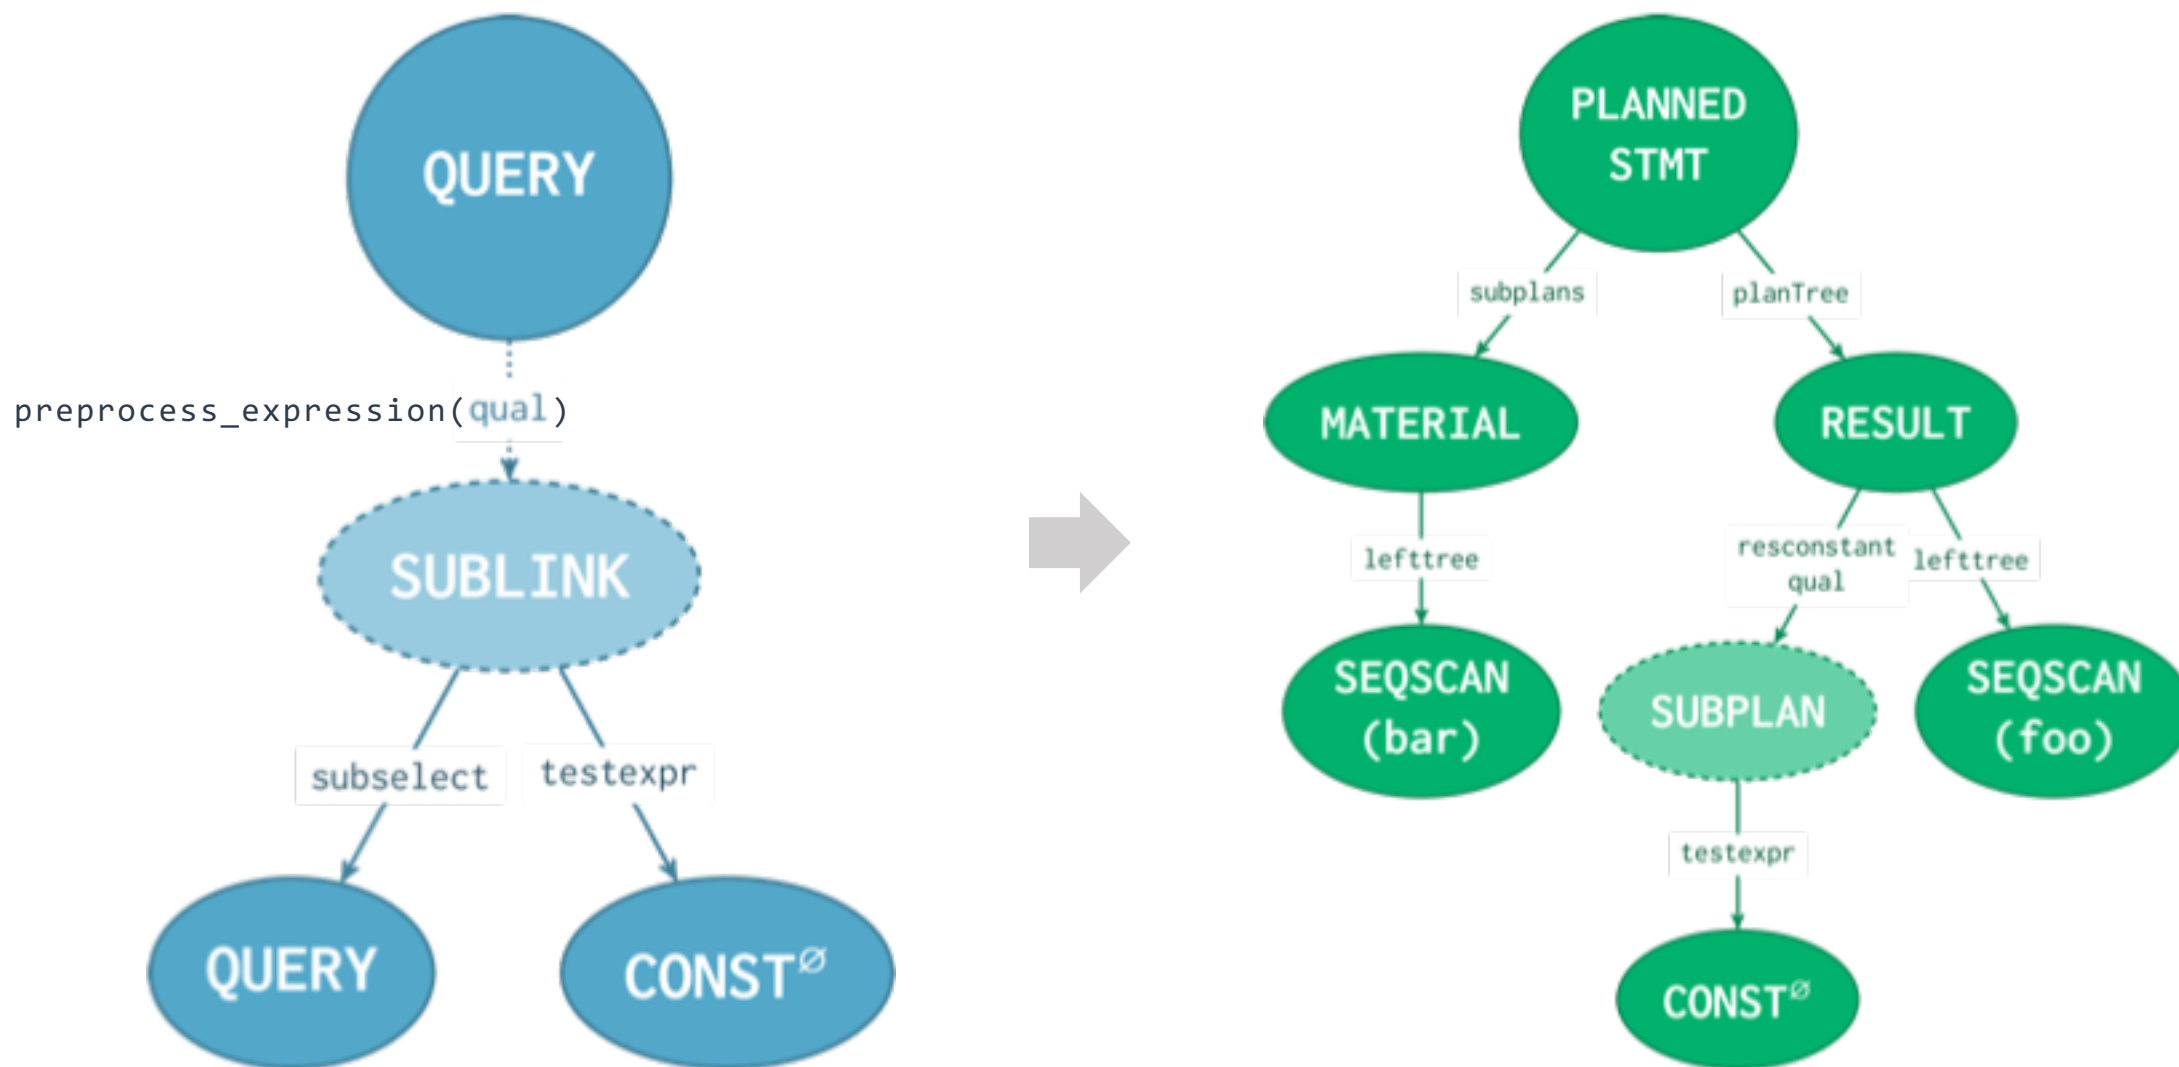
ANY Sublink Pullup

SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);

Current Pre-processed Query Tree


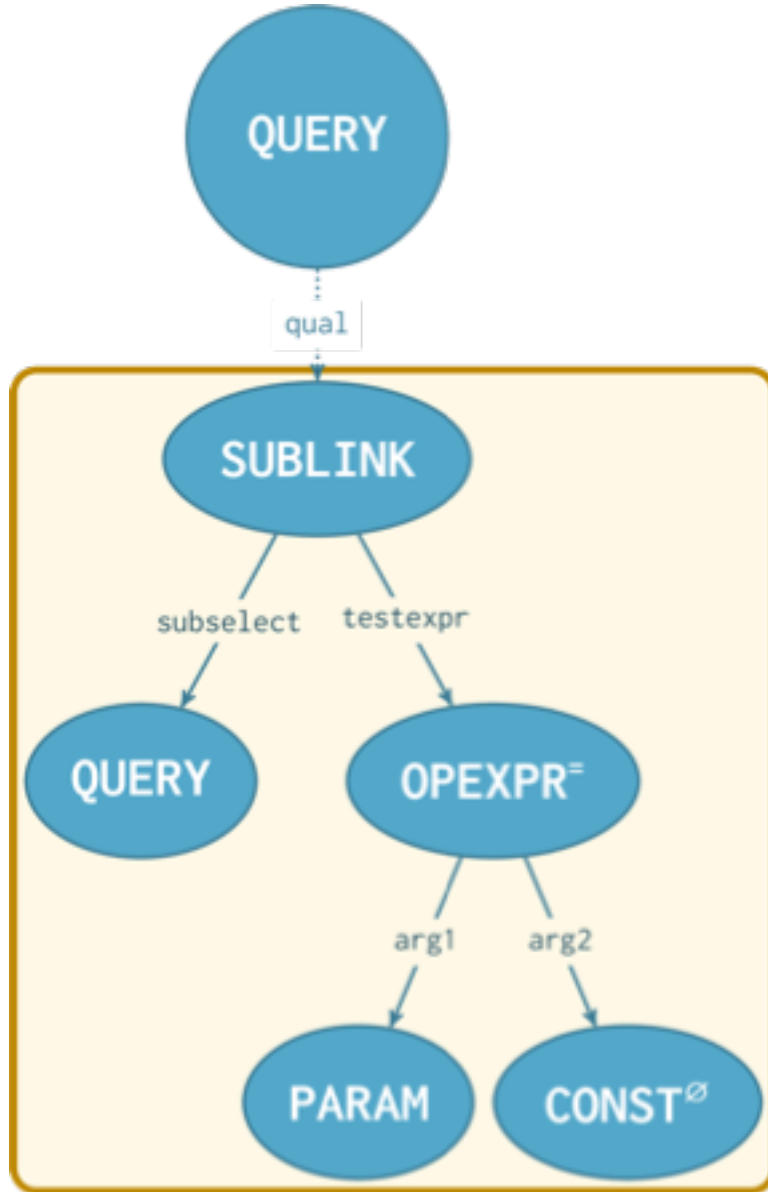


SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);




```
# SELECT a FROM foo WHERE NULL  
= ANY(SELECT b FROM bar);
```

```
# SELECT a FROM foo WHERE NULL;
```



~~Rule 1~~

This is semantically incorrect in one case

NULL Semantics

Meet ANY semantics

NULL $\stackrel{?}{=}$ ANY(SELECT b FROM bar)

Does any b in bar equal an unknown?

```
# SELECT NULL = ANY(SELECT b FROM bar);
```

Does any *b* in *bar* equal an unknown?

Does any **b** in **bar** equal an **unknown**?

```
# SELECT NULL = ANY(SELECT  
b FROM bar);
```

?column?

(1 row)

Does any **b** in **bar** equal an **unknown**?

```
# SELECT NULL = ANY(SELECT  
b FROM bar);
```

?column?

(1 row)

```
# TRUNCATE bar;
```

```
# SELECT NULL = ANY(SELECT  
b FROM bar);
```

?column?

f

(1 row)

```
# SELECT a FROM foo
  WHERE NULL = ANY(
    SELECT b FROM bar
  );
```

a

(0 rows)

```
# TRUNCATE bar;
```

```
# SELECT a FROM foo
  WHERE NULL = ANY(
    SELECT b FROM bar
  );
```

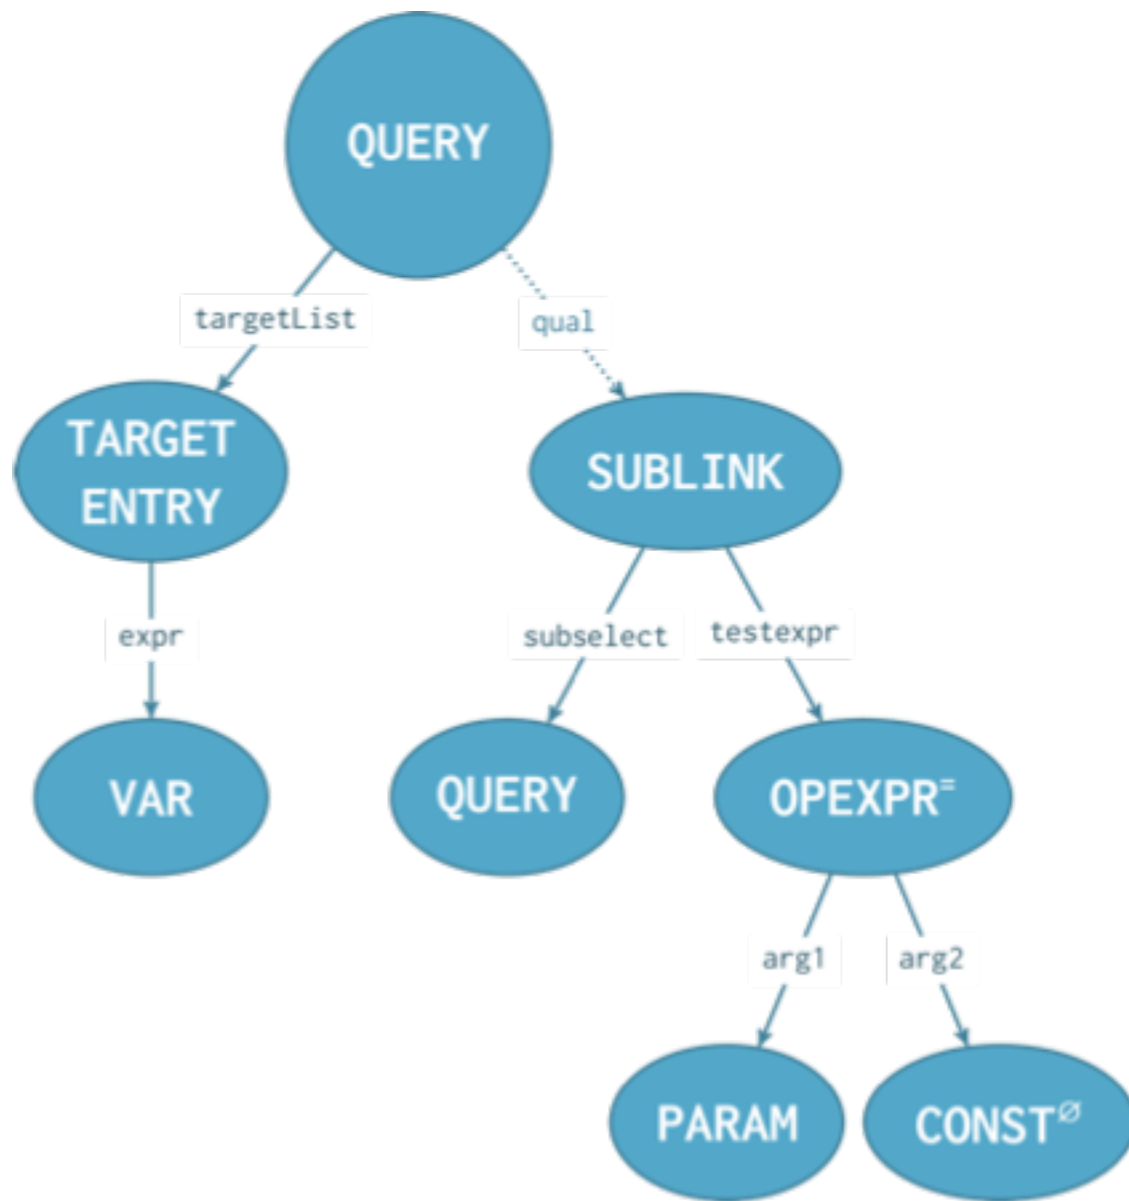
a

(0 rows)


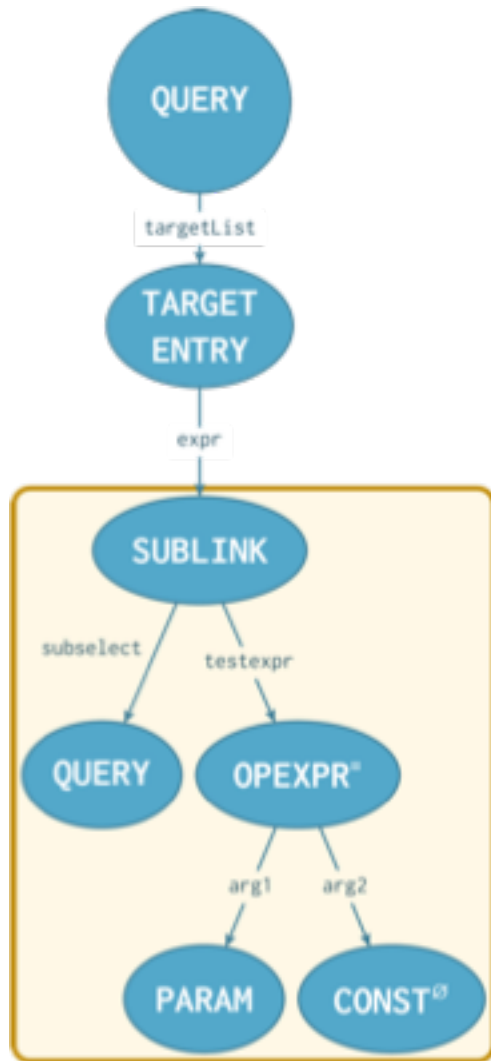
SELECT NULL = ANY(SELECT b FROM bar);



SELECT a FROM foo WHERE
NULL = ANY(SELECT b FROM bar);



FALSE if **bar** is an empty table and **NULL** otherwise



What could we do instead?

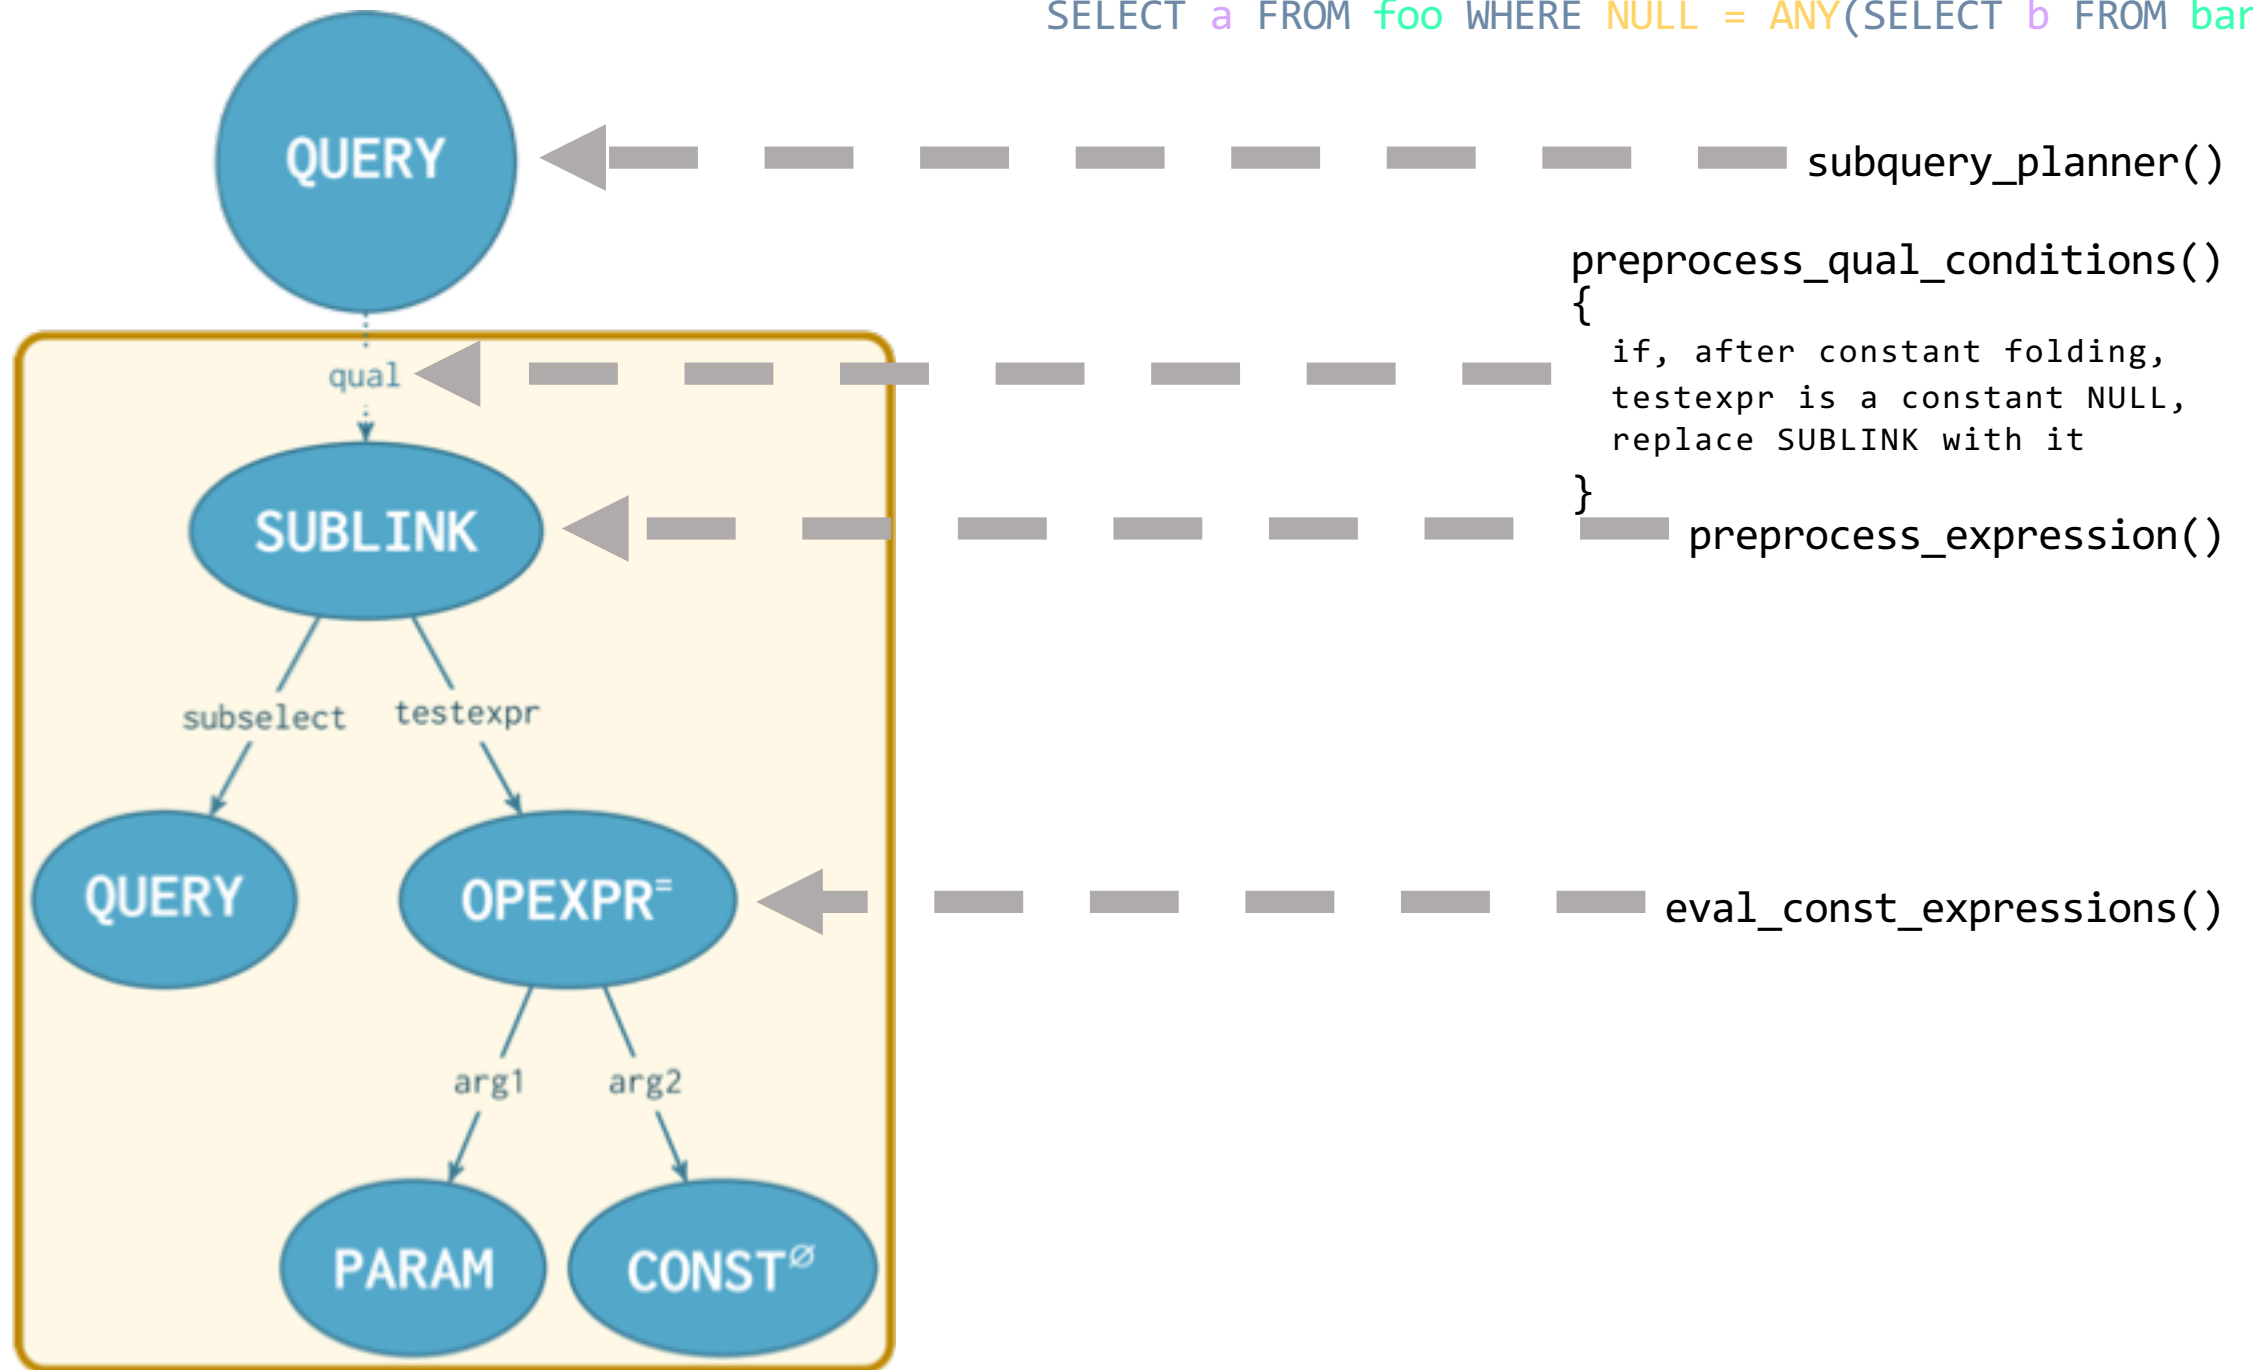
Two s

Constant Folding only in the qual

ANY Sublink Pullup


```
# SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

```
SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```



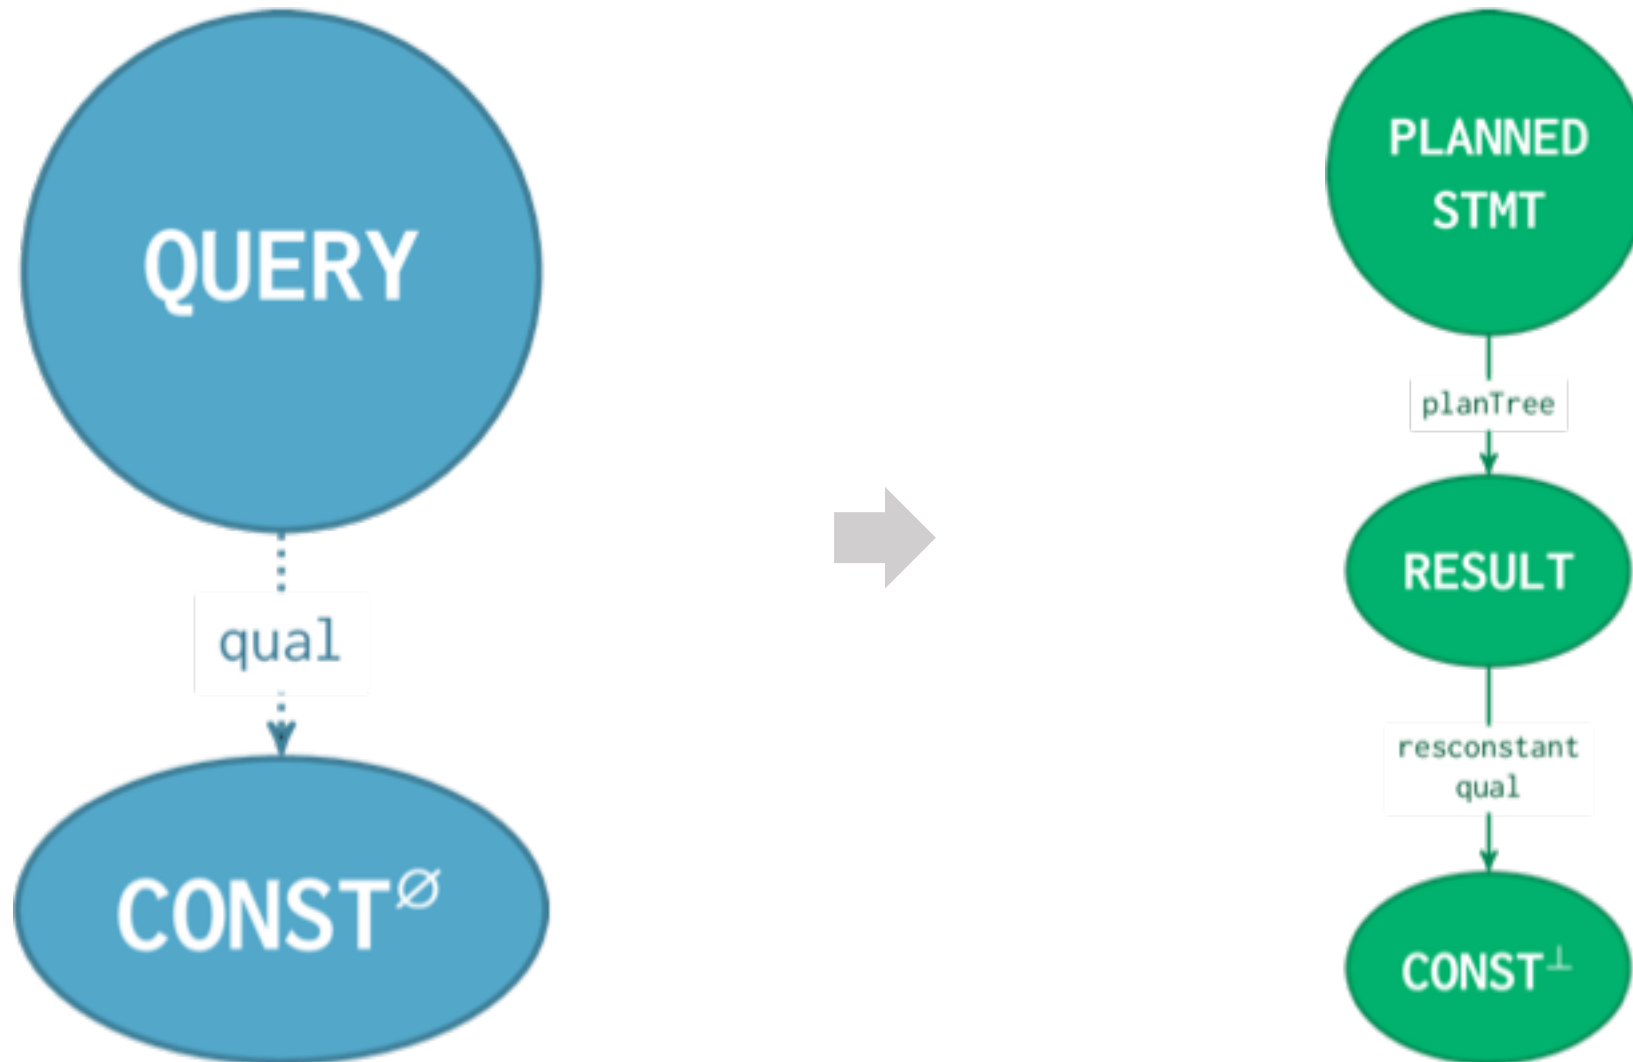
SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);

Replace **ANY SUBLINK** when pre-processing **quals**



```
SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

Patched Planning



Patched Plan

```
# EXPLAIN SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

QUERY PLAN

Result (cost=... rows=0 width=...)

One-Time Filter: false

~~Rule 4~~

A very narrow case

Two  s

Constant Folding

ANY Sublink Pullup

```
# SELECT a FROM foo WHERE NULL = ANY(SELECT b FROM bar);
```

EXPLAIN SELECT a FROM foo WHERE a = ANY(SELECT b FROM bar);

QUERY PLAN

Hash Join

Hash Cond: (foo.a = bar.b)

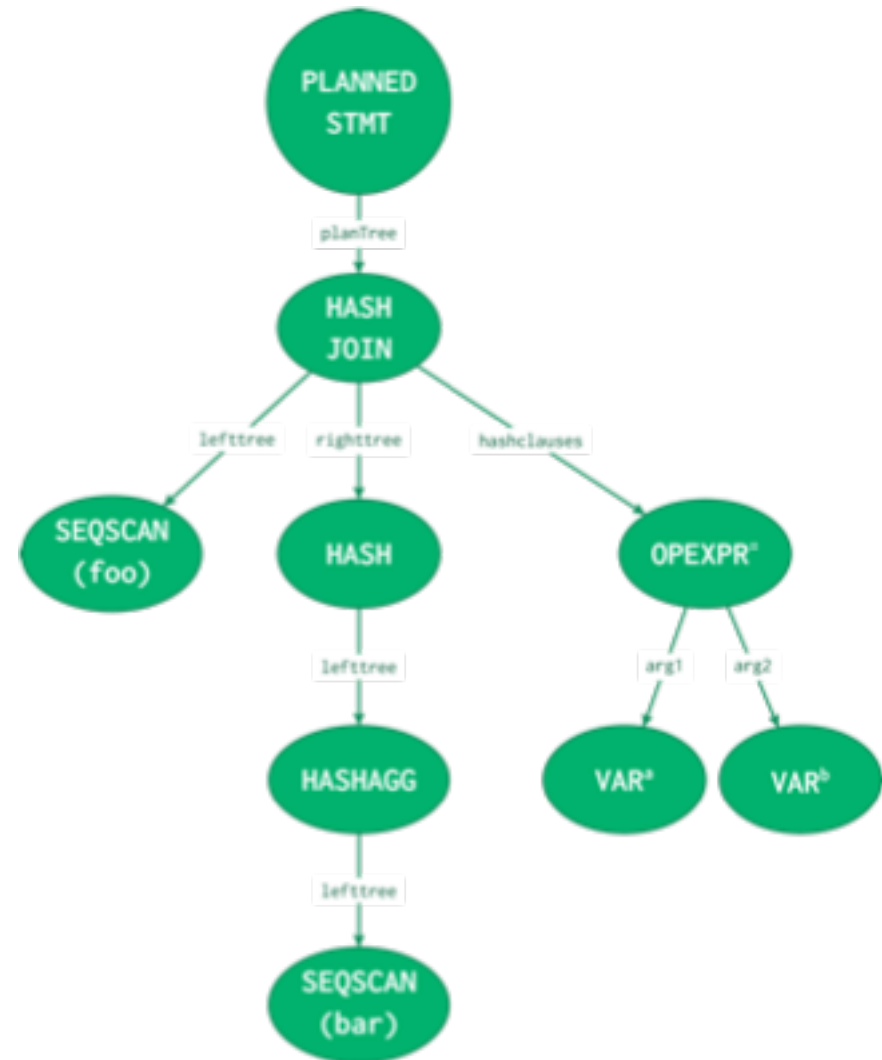
→ Seq Scan on foo

→ Hash

→ HashAggregate


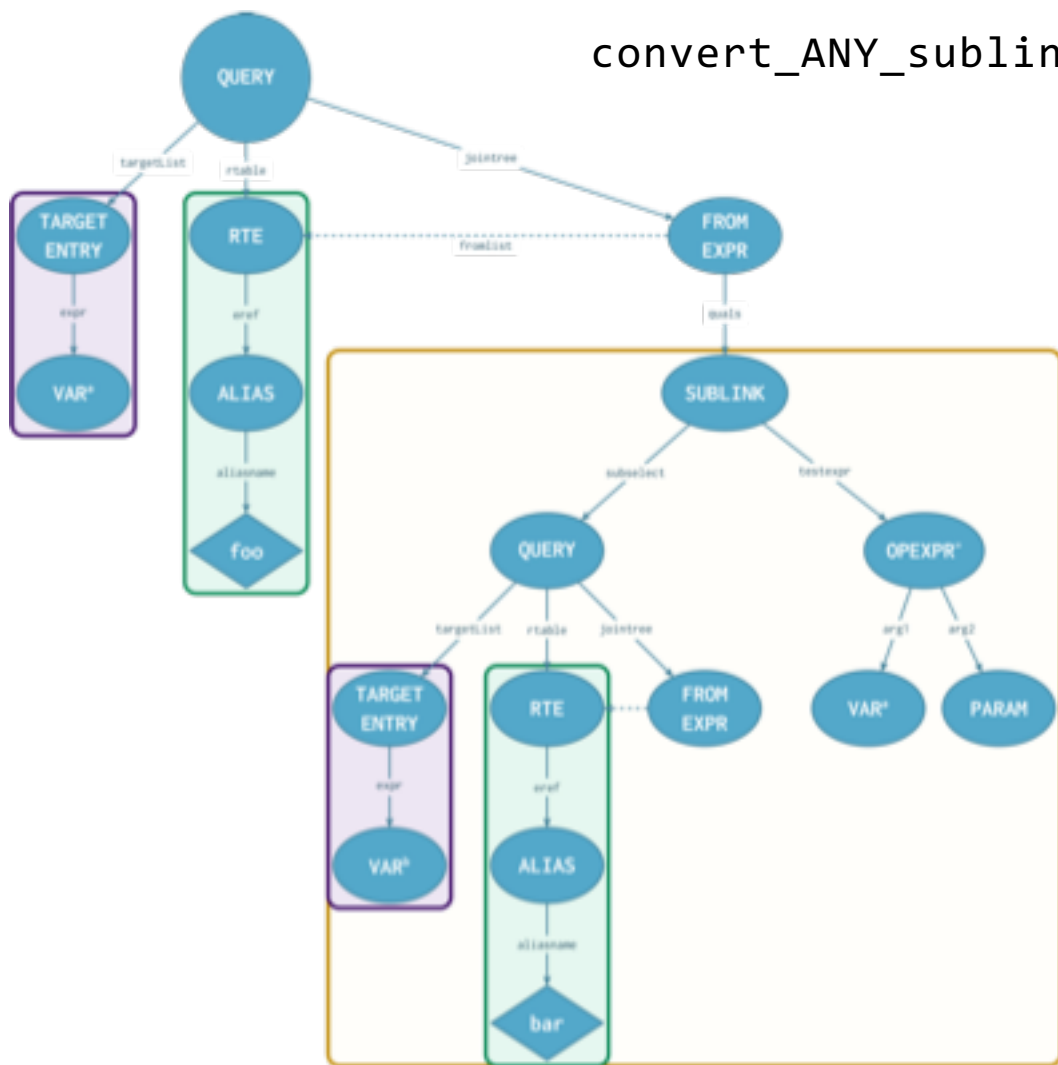
Group Key: bar.b

→ Seq Scan on bar




SELECT **a** FROM **foo** WHERE **a** = ANY(SELECT **b** FROM **bar**);


convert_ANY_sublink_to_join()




... **a** = **ANY**(SELECT **b** FROM **bar**);




... **NULL** = **ANY**(SELECT **b** FROM **bar**);



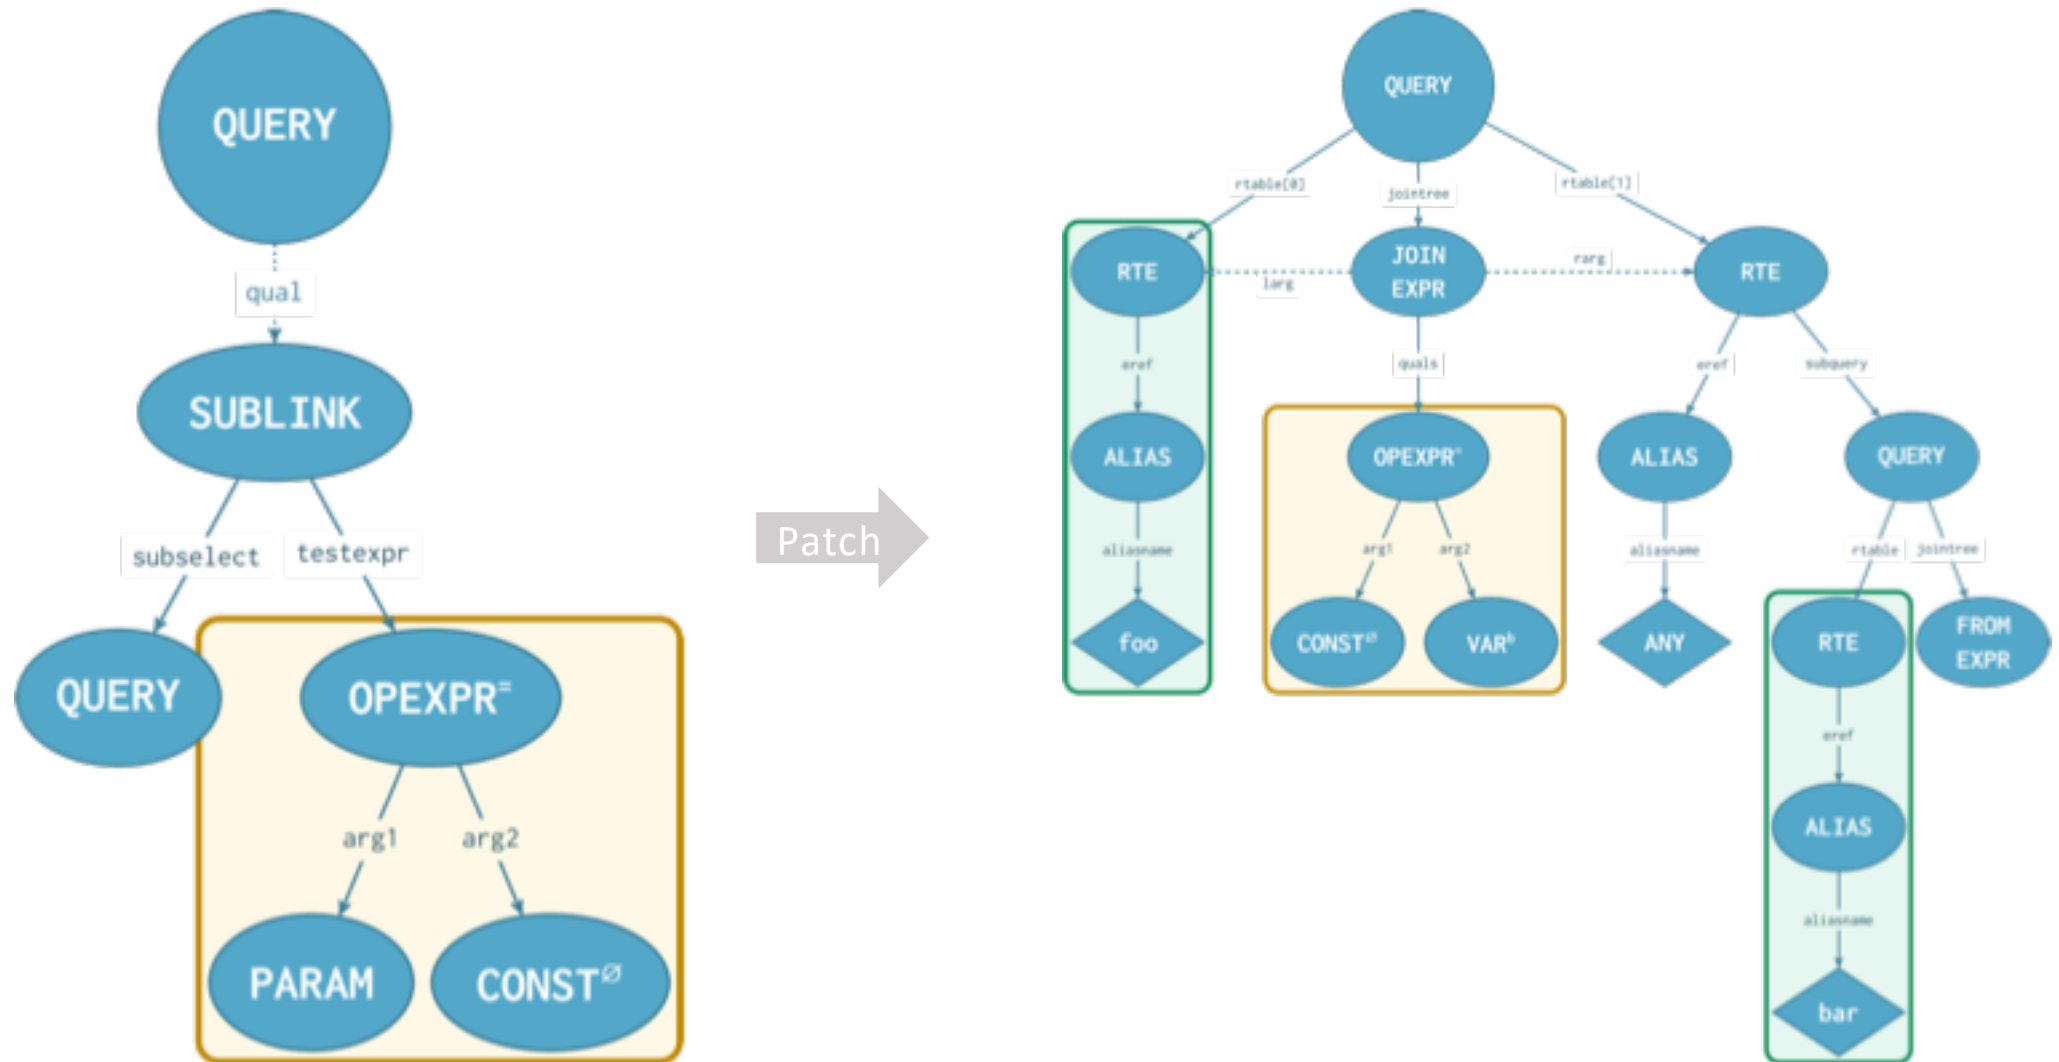
... **a** = **ANY**(SELECT **b** FROM **bar**);



... **NULL** = **ANY**(SELECT **b** FROM **bar**);


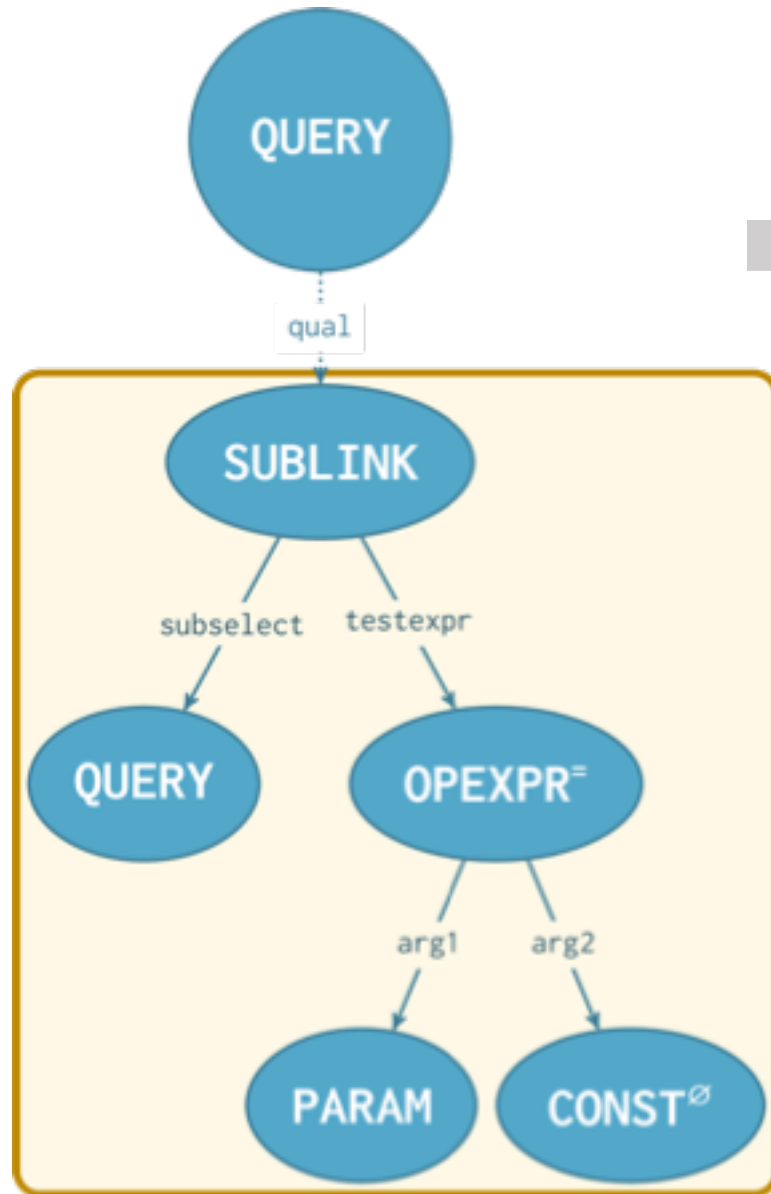


SELECT **a** FROM **foo** WHERE **NULL** = **ANY**(SELECT **b** FROM **bar**);



SELECT a FROM foo WHERE NULL
= ANY(SELECT b FROM bar);

SELECT a FROM foo JOIN bar
WHERE NULL = deduped(b);



```
# EXPLAIN SELECT a FROM foo WHERE 7 = ANY(SELECT b FROM bar WHERE b = 5);
```

Current

QUERY PLAN

Result

One-Time Filter: (hashed SubPlan 1)
→ Seq Scan on foo
SubPlan 1
→ Seq Scan on bar
Filter: (b = 7)

Patched

QUERY PLAN

Result

One-Time Filter: false

```
# EXPLAIN SELECT a FROM foo WHERE 7 = ANY(SELECT b FROM bar);
```

Current

QUERY PLAN

Result

One-Time Filter: (hashed SubPlan 1)
→ Seq Scan on foo
SubPlan 1
→ Seq Scan on bar

Patched

QUERY PLAN

Nested Loop Semi Join

→ Seq Scan on foo
→ Materialize
→ Seq Scan on bar
Filter: (7 = b)

~~Rules ②, ④~~

Produces worse plans when the join isn't eliminated

A very narrow case

Guidelines for New Optimizations

- ① Does it always retain semantic correctness?
- ② Does it inhibit downstream optimizations?
- ③ Is the improvement in execution time worth the cost in planning time?
- ④ Is the complexity cost commensurate with the performance benefit?

Some Rejected s

- Use stats
- Execute the subquery



Discussion



When is it okay to ...?

- Do a catalog lookup
- Do partial execution
- Mutate the plan tree
- Save a reference to parent query

Guidelines ... Others?

- ① Does it always retain semantic correctness?
- ② Does it inhibit downstream optimizations?
- ③ Is the improvement in execution time worth the cost in planning time?
- ④ Is the complexity cost commensurate with the performance benefit?

(Re)sources

- Uncommitted planner patches and discussion (browse old commitfests) <https://commitfest.postgresql.org/>
- Planner hacking presentations
 - Tom Lane PGCon 2011 Hacking the Query Planner
https://www.pgcon.org/2011/schedule/attachments/188_Planner%20talk.pdf
 - Robert Haas (CTRL-F 'planner')
<https://sites.google.com/site/robertmhaas/presentations/2010-2012>
- `src/backend/optimizer/README`

github.com/melanieplageman/

↳ /debugging_planner

Slides and Glossary

↳ /postgres/tree/

Code

↳ /const_folding_sublink_wrong

Constant Folding

↳ /qual_scoped_const_folding_sublink

Constant Folding only in the qual

↳ /const_ANY_sublink_pullup

ANY Sublink Pullup

Acknowledgements

Jesse Zhang – Queries and content assistance

Kaiting Chen—TikZ diagram designer